

Week 3: CNN Cancer Detection Kaggle Mini-Project

Problem Description:

Build a deep learning model to perform binary image classification on image patches to identify metastatic cancer.

1. Import the necessary python libraries

```
import os
import gc
from pathlib import Path
from math import ceil
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import tensorflow as tf
from tensorflow.keras import layers, models
from keras.preprocessing.image import load_img, img_to_array
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
```

2. Load Data

Get the path to the test and train TIFFs and read in the csv train_labels.csv and sample_submission.csv data. The first 5 rows of the train_labels.csv are displayed as well.

```
test_data_path = os.path.join(os.getcwd(), "histopathologic-cancer-
detection/test")
train_data_path = os.path.join(os.getcwd(), "histopathologic-cancer-
detection/train")

train_label_path = os.path.join(os.getcwd(), "histopathologic-cancer-
detection/train_labels.csv")
train_labels = pd.read_csv(train_label_path)
display(train_labels.head())
```

	id	label
0	f38a6374c348f90b587e046aac6079959adf3835	0
1	c18f2d887b7ae4f6742ee445113fa1aef383ed77	1
2	755db6279dae599ebb4d39a9123cce439965282d	0
3	bc3f0c64fb968ff4a8bd33af6971ecae77c75e08	0
4	068aba587a4950175d04c680d38943fd488d6a9d	0

3. Exploratory Data Analysis (EDA)

Next perform some Exploratory Data Analysis. For an initial data analysis, check for any null values in the training dataset, by calling `isna()` and `sum()` on the training labels dataframe. The return of the function displays 0 for both columns so looks like there are no null values/data in the dataset.

```
display(train_labels.isna().sum())  
  
id      0  
label   0  
dtype: int64
```

Next check for basic info about the training dataset, which shows that the dataset only has two columns "id", and "label" and 220025 rows of data.

```
train_labels.info()  
  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 220025 entries, 0 to 220024  
Data columns (total 2 columns):  
#   Column  Non-Null Count  Dtype  
---  -----  -  
0    id      220025 non-null   object  
1    label    220025 non-null   int64  
dtypes: int64(1), object(1)  
memory usage: 3.4+ MB
```

Another form of EDA that can be performed is the `describe()` functions that will display some basic metrics/statistics. Some of these metrics include total row counts, mean, and standard deviation.

```
display(train_labels.describe())  
  
          label  
count  220025.000000  
mean         0.405031  
std          0.490899  
min          0.000000  
25%          0.000000  
50%          0.000000  
75%          1.000000  
max          1.000000
```

Next as part of the EDA calculate the label distribution, and create a bar plot to visually display the calculated distribution. The label percentage is also calculated and displayed. Based off the label distribution bar plot and the label percentage the labels are for the most part balanced. There is a slight skew of label 0 which is about 59.5% percent but it is not too high enough to be a problem during training.


```

label_counts = train_labels["label"].value_counts().sort_index()

plt.figure()
label_counts.plot(kind="bar")
plt.title("Label Distribution")
plt.xlabel("Label\n(0 = No Tumor, 1 = Tumor)")
plt.ylabel("Count")
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()

label_percentage = (label_counts / label_counts.sum()) * 100
print(f"Label Percentage:")
display(label_percentage)

```



Label Percentage:

```

label
0    59.496875
1    40.503125
Name: count, dtype: float64

```

Next as part of the EDA plot a subset of the TIFFs images from the the training set and test set. This will help visully see what the data looks.

```
file_paths = []

for id in train_labels.head(9)["id"]:
    full_path = os.path.join(train_data_path, f"{id}.tif")
    file_paths.append(full_path)

n_imgs = len(file_paths)
cols = 3
rows = ceil(n_imgs / cols)
figsize = (4 * cols, 4 * rows)

fig, axes = plt.subplots(rows, cols, figsize=figsize)
fig.suptitle("Subsample of Training Images", fontsize=18)
axes = axes.ravel()


for idx, (ax, path) in enumerate(zip(axes, file_paths)):
    img = load_img(path)
    ax.imshow(img)
    ax.set_title(Path(path).stem, fontsize=10)
    ax.axis("off")

for ax in axes[n_imgs:]:
    ax.axis("off")


plt.tight_layout()
plt.show()
```

Subsample of Training Images


f38a6374c348f90b587e046aac6079959adf3835




c18f2d887b7ae4f6742ee445113fa1aef383ed77




755db6279dae599ebb4d39a9123cce439965282d




bc3f0c64fb968ff4a8bd33af6971ecae77c75e08




068aba587a4950175d04c680d38943fd488d6a9d




acfe80838488fae3c89bd21ade75be5c34e66be7




a24ce148f6ffa7ef8eeb4efb12ebffe8dd700da



7f6ccae485af121e0b6ee733022e226ee6b0c65f



559e55a64c9ba828f700e948f6886f4cea919261



```
tiffs = ["0a5ebbbbba69285b622762de2bd790e06d6862f1c.tif",
"0a0a85db9218e366569c913185cc0740f59f4d9e.tif",
"0a0c08f8f3a830f5d925027e90a889b9c00cf518.tif",
"0a0cfd6408ef2e4f94618c8a37b16d47ed840468.tif",
"0a0d0532a5578a852f5c6b831cce998cld9b450c.tif",
"0a0e1de0d1bd38eff2942efc13e5e0420597f295.tif",
"0a0e554e88e535e0b57cc5a60ee0458802b90dfb.tif",
"0a0f66c67c6583cab23bbb7d693a824d5c54f8ce.tif",
"00a01a16ea56bcc9463351b6a5c3ca0fb0bf114c.tif"
]
```

```
file_paths = []

for tiff in tiffs:
    full_path = os.path.join(test_data_path, tiff)
    file_paths.append(full_path)

#display(file_paths)


n_imgs = len(file_paths)
cols = 3
rows = ceil(n_imgs / cols)
figsize = (4 * cols, 4 * rows)

fig, axes = plt.subplots(rows, cols, figsize=figsize)
fig.suptitle("Subsample of Test Images", fontsize=18)
axes = axes.ravel()

for idx, (ax, path) in enumerate(zip(axes, file_paths)):
    img = load_img(path)
    ax.imshow(img)
    ax.set_title(Path(path).stem, fontsize=10)
    ax.axis("off")

for ax in axes[n_imgs:]:
    ax.axis("off")

plt.tight_layout()
plt.show()
```



4. Prepare data from model training.

Since the Exploratory Data Analysis did not reveal and hidden bias or missing data in the dataset. There is not much data cleaning/preprocessing necessary. However, the dataset does need to be prepared before training a model can begin. For this a helper function is created to prepare the dataset. The `prepare_dataset()` function will iterate through all labels in the dataset and build the image path directory to load the specific TIF image which is then append a list along with the label in a separate list. The lists are then converted to a numpy array and returned.

```
def prepare_dataset(dataset:pd.DataFrame, image_dir:str, ):
    x = []
```

```

y = []
for idx, row in dataset.iterrows():
    img_path = os.path.join(image_dir, f"{row['id']}.tif")
    img = load_img(img_path, target_size=(64, 64))
    img_array = img_to_array(img) / 255.0
    x.append(img_array)
    y.append(row["label"])

return np.array(x), np.array(y)

x, y = prepare_dataset(train_labels, train_data_path)

```

The next step in preparing the dataset for training a model is splitting the dataset into a training and testing subset of data. This is achieved with the function `train_test_split()` which will split the main dataset into 80% training and 20% testing.

```

x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, random_state=42)

```

4. Model 1 Architecture

For the first model I decided to keep it simple and build a CNN model using TensorFlow's Keras API. It consists of the following layers:

1. Input Layer
 - Accepts images of size (64, 64, 3) corresponding to 64x64 pixel color image (RGB)
2. First convolutional Block
 - Conv2D Layer: 32 filters, 3x3 kernel size, ReLU activation.
 - MaxPooling2D Layer: 2x2 pool size.
 - This allows the layer to extract basic low-level features from the images, such as edges, corners, and simple textures.
3. Second Convolutional Block
 - Conv2D Layer: 64 filters, 3x3 kernel size, ReLU activation.
 - MaxPooling2D Layer: 2x2 pool size.
 - Builds on the first layer to detect more complex patterns and textures by combining simple features.
4. Flatten Layer
 - Flattens the 2D feature maps into a 1D vector for the fully connected layers.
5. Fully Connected
 - Dense Layer: 128 units, ReLU activation.
 - Learns complex combinations of the extracted features to form high-level representations important for classification


```

image_size = 64

cnn_model_1 = models.Sequential([
    layers.Conv2D(32, (3, 3), activation="relu",
input_shape=(image_size, image_size, 3)),
    layers.MaxPooling2D((2,2)),

    layers.Conv2D(64, (3, 3), activation="relu"),
    layers.MaxPooling2D((2,2)),

    layers.Flatten(),
    layers.Dense(128, activation="relu"),
    layers.Dropout(0.5),
    layers.Dense(1, activation="sigmoid")
])

cnn_model_1.compile(optimizer="adam", loss="binary_crossentropy",
metrics=["accuracy"])

cnn_model_1.summary()

c:\Users\charlie\.envs\deep_learning_3_11\Lib\site-packages\keras\src\
layers\convolutional\base_conv.py:107: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in
the model instead.
  super().__init__(activity_regularizer=activity_regularizer,
**kwargs)

```

Model: "sequential"

Layer (type) Param #	Output Shape	
conv2d (Conv2D) 896	(None, 62, 62, 32)	
max_pooling2d (MaxPooling2D) 0	(None, 31, 31, 32)	
conv2d_1 (Conv2D) 18,496	(None, 29, 29, 64)	
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 64)	

0				
		flatten (Flatten)	(None, 12544)	
0				
		dense (Dense)	(None, 128)	
1,605,760				
		dropout (Dropout)	(None, 128)	
0				
		dense_1 (Dense)	(None, 1)	
129				
Total params: 1,625,281 (6.20 MB)				
Trainable params: 1,625,281 (6.20 MB)				
Non-trainable params: 0 (0.00 B)				

After creating the CNN model call the fit() method on it to perform the training on the model. The training will use the training subset and train over 20 epochs with a batch_size of 32 and with a validation split of 20%

```
trained_cnn_model_1 = cnn_model_1.fit(x_train, y_train, epochs=20,
batch_size=32, validation_split=0.2)
```

Epoch 1/20
4401/4401 ————— 29s 6ms/step - accuracy: 0.7506 - loss: 0.5033 - val_accuracy: 0.8268 - val_loss: 0.3968

Epoch 2/20
4401/4401 ————— 23s 5ms/step - accuracy: 0.8299 - loss: 0.3931 - val_accuracy: 0.8500 - val_loss: 0.3566

Epoch 3/20
4401/4401 ————— 23s 5ms/step - accuracy: 0.8455 - loss: 0.3590 - val_accuracy: 0.8501 - val_loss: 0.3469

Epoch 4/20
4401/4401 ————— 21s 5ms/step - accuracy: 0.8509 - loss: 0.3470 - val_accuracy: 0.8521 - val_loss: 0.3366

Epoch 5/20
4401/4401 ————— 21s 5ms/step - accuracy: 0.8551 - loss: 0.3369 - val_accuracy: 0.8436 - val_loss: 0.3933

Epoch 6/20
4401/4401 ————— 22s 5ms/step - accuracy: 0.8590 - loss:

```

0.3281 - val_accuracy: 0.8576 - val_loss: 0.3364
Epoch 7/20
4401/4401 _____ 21s 5ms/step - accuracy: 0.8660 - loss:
0.3159 - val_accuracy: 0.8607 - val_loss: 0.3311
Epoch 8/20
4401/4401 _____ 22s 5ms/step - accuracy: 0.8662 - loss:
0.3124 - val_accuracy: 0.8554 - val_loss: 0.3346
Epoch 9/20
4401/4401 _____ 23s 5ms/step - accuracy: 0.8697 - loss:
0.3062 - val_accuracy: 0.8557 - val_loss: 0.3408
Epoch 10/20
4401/4401 _____ 23s 5ms/step - accuracy: 0.8773 - loss:
0.2936 - val_accuracy: 0.8594 - val_loss: 0.3322
Epoch 11/20
4401/4401 _____ 23s 5ms/step - accuracy: 0.8801 - loss:
0.2883 - val_accuracy: 0.8576 - val_loss: 0.3382
Epoch 12/20
4401/4401 _____ 22s 5ms/step - accuracy: 0.8834 - loss:
0.2779 - val_accuracy: 0.8609 - val_loss: 0.3352
Epoch 13/20
4401/4401 _____ 23s 5ms/step - accuracy: 0.8885 - loss:
0.2695 - val_accuracy: 0.8609 - val_loss: 0.3392
Epoch 14/20
4401/4401 _____ 23s 5ms/step - accuracy: 0.8904 - loss:
0.2644 - val_accuracy: 0.8540 - val_loss: 0.3750
Epoch 15/20
4401/4401 _____ 22s 5ms/step - accuracy: 0.8961 - loss:
0.2541 - val_accuracy: 0.8648 - val_loss: 0.3387
Epoch 16/20
4401/4401 _____ 23s 5ms/step - accuracy: 0.9005 - loss:
0.2420 - val_accuracy: 0.8624 - val_loss: 0.3620
Epoch 17/20
4401/4401 _____ 22s 5ms/step - accuracy: 0.9038 - loss:
0.2362 - val_accuracy: 0.8642 - val_loss: 0.3621
Epoch 18/20
4401/4401 _____ 22s 5ms/step - accuracy: 0.9075 - loss:
0.2262 - val_accuracy: 0.8619 - val_loss: 0.3755
Epoch 19/20
4401/4401 _____ 24s 5ms/step - accuracy: 0.9126 - loss:
0.2138 - val_accuracy: 0.8574 - val_loss: 0.3834
Epoch 20/20
4401/4401 _____ 22s 5ms/step - accuracy: 0.9156 - loss:
0.2111 - val_accuracy: 0.8626 - val_loss: 0.3637

```

After training the model the next step is to evaluate the performance of the trained model by using the test subset to score the accuracy of the model. The accuracy of the model resulted a accuracy score of about 86.7%. Which is a relatively high score.


```
cnn_model_1_test_loss, cnn_model_1_test_acc =  
cnn_model_1.evaluate(x_test, y_test, verbose=2)
```

```
print(f"Test Acc: {cnn_model_1_test_acc:.3f}")
```

```
1376/1376 - 2s - 1ms/step - accuracy: 0.8643 - loss: 0.3540  
Test Acc: 0.864
```

Next a plot of the training and validation accuracy over each epoch is plotted. This to aid in visualizing how the model learned over time. Based of the plot the model has good genralizaion since the validation accuracy raises with the training accuracy.

```
plt.plot(trained_cnn_model_1.history["accuracy"], label="Train Acc")  
plt.plot(trained_cnn_model_1.history["val_accuracy"], label="Val Acc")  
plt.xlabel("Epoch")  
plt.ylabel("Accuracy")  
plt.legend()  
plt.title("Model 1: Training vs Validation Accuracy")  
plt.show()
```



Model 1 Tuning

In an effort to improve the accuracy of the model, some tuning will be attempted. The following parameters were tuned in order to increase performance:

- Increased the number filters doubling the capacity at each convolutional layer.
- Reduced dropout rate from 50% to 40% to keep more information while still controlling overfitting.
- Lowered the learning rate to make smaller, more precise weight updates from 0.001 -> 0.005.
- Added EarlyStopping to automatically stop training if validation loss does not improve, preventing overfitting and saving training time.

```
cnn_model_1_tuned = models.Sequential([
    layers.Conv2D(64, (3, 3), activation="relu",
input_shape=(image_size, image_size, 3)),
    layers.MaxPooling2D((2,2)),

    layers.Conv2D(128, (3, 3), activation="relu"),
    layers.MaxPooling2D((2,2)),

    layers.Flatten(),
    layers.Dense(128, activation="relu"),
    layers.Dropout(0.4),
    layers.Dense(1, activation="sigmoid")
])
```

```
cnn_model_1_tuned.compile(optimizer=Adam(learning_rate=0.0005),
loss="binary_crossentropy", metrics=["accuracy"])
```

```
early_stop = EarlyStopping(monitor="val_loss", patience=5,
restore_best_weights=True)
```

```
cnn_model_1_tuned.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	
Param #		
conv2d_2 (Conv2D)	(None, 62, 62, 64)	
1,792		
max_pooling2d_2 (MaxPooling2D)	(None, 31, 31, 64)	
0		

conv2d_3 (Conv2D)	(None, 29, 29, 128)	
73,856		
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 128)	
0		
flatten_1 (Flatten)	(None, 25088)	
0		
dense_2 (Dense)	(None, 128)	
3,211,392		
dropout_1 (Dropout)	(None, 128)	
0		
dense_3 (Dense)	(None, 1)	
129		

Total params: 3,287,169 (12.54 MB)

Trainable params: 3,287,169 (12.54 MB)

Non-trainable params: 0 (0.00 B)

After tuning the new CNN model call the fit() method on it to perform the training on the model. The training will use the training subset and train over 20 epochs with a batch_size of 32 and with a validation split of 20%

```
trained_cnn_model_1_tuned = cnn_model_1_tuned.fit(x_train, y_train,
epochs=30, batch_size=64, validation_split=0.2,
callbacks=[early_stop], verbose=2)
```

Epoch 1/30

2201/2201 - 64s - 29ms/step - accuracy: 0.7885 - loss: 0.4646 -
val_accuracy: 0.8044 - val_loss: 0.4303

Epoch 2/30

2201/2201 - 63s - 29ms/step - accuracy: 0.8272 - loss: 0.3991 -
val_accuracy: 0.8325 - val_loss: 0.3710

Epoch 3/30

2201/2201 - 64s - 29ms/step - accuracy: 0.8495 - loss: 0.3530 -
val_accuracy: 0.8596 - val_loss: 0.3284

Epoch 4/30

2201/2201 - 64s - 29ms/step - accuracy: 0.8611 - loss: 0.3275 -

```

val_accuracy: 0.8551 - val_loss: 0.3302
Epoch 5/30
2201/2201 - 64s - 29ms/step - accuracy: 0.8718 - loss: 0.3034 -
val_accuracy: 0.8591 - val_loss: 0.3369
Epoch 6/30
2201/2201 - 63s - 28ms/step - accuracy: 0.8828 - loss: 0.2809 -
val_accuracy: 0.8651 - val_loss: 0.3276
Epoch 7/30
2201/2201 - 63s - 29ms/step - accuracy: 0.8925 - loss: 0.2595 -
val_accuracy: 0.8722 - val_loss: 0.3014
Epoch 8/30
2201/2201 - 63s - 28ms/step - accuracy: 0.9035 - loss: 0.2369 -
val_accuracy: 0.8698 - val_loss: 0.3051
Epoch 9/30
2201/2201 - 62s - 28ms/step - accuracy: 0.9140 - loss: 0.2128 -
val_accuracy: 0.8741 - val_loss: 0.3102
Epoch 10/30
2201/2201 - 63s - 28ms/step - accuracy: 0.9268 - loss: 0.1867 -
val_accuracy: 0.8657 - val_loss: 0.3451
Epoch 11/30
2201/2201 - 63s - 29ms/step - accuracy: 0.9377 - loss: 0.1621 -
val_accuracy: 0.8682 - val_loss: 0.3589
Epoch 12/30
2201/2201 - 64s - 29ms/step - accuracy: 0.9480 - loss: 0.1387 -
val_accuracy: 0.8703 - val_loss: 0.4076

```

Re-evaluate the accuracy on the testing subset to check if the tuning helped increasing the accuracy of the model. Based off the accuracy score of 87.9 there was a slightly improvement in the model after the tuning although not by much.

```

cnn_model_1_tuned_test_loss, cnn_model_1_tuned_test_acc =
cnn_model_1_tuned.evaluate(x_test, y_test, verbose=2)

print(f"Test Acc: {cnn_model_1_tuned_test_acc:.3f}")

1376/1376 - 4s - 3ms/step - accuracy: 0.8755 - loss: 0.2958
Test Acc: 0.875


```

Plot the Training vs Validation Accuracy plot.

```

plt.plot(trained_cnn_model_1_tuned.history["accuracy"], label="Train
Acc")
plt.plot(trained_cnn_model_1_tuned.history["val_accuracy"], label="Val
Acc")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()
plt.title("Model 1 Tuned: Training vs Validation Accuracy")
plt.show()

```



5. Model 2 Architecture

For the second model I decided to design a slightly deeper CNN architecture using TensorFlow's Keras API. It consists of the following layers:

1. Input Layer:
 - Accepts images of (64, 64, 3), corresponding to 64x64 pixel color images (RGB).
2. First Convolutional Block:
 - Conv2D Layer: 16 filters, 3x3 kernel size, ReLU activation, "same" padding.
 - Conv2D Layer: 16 filters, 3x3 kernel size, ReLU activation, "same" padding.
 - MaxPooling2D Layer: 2x2 pool size.
 - This block allows the model to extract simple low-level features such as edges and textures, while maintaining the spatial dimensions due to "same" padding.
3. Second Convolutional Block:
 - Conv2D Layer: 32 filters, 3x3 kernel size, ReLU activation, "same" padding.
 - Conv2D Layer: 32 filters, 3x3 kernel size, ReLU activation, "same" padding.
 - MaxPooling2D Layer: 2x2 pool size.
 - This block builds on the first, enabling the model to detect more detailed features and patterns by combining simpler features.

4. Third Convolutional Block:

- Conv2D Layer: 64 filters, 3x3 kernel size, ReLU activation, "same" padding.
- MaxPooling2D Layer: 2x2 pool size.
- This layer further deepens the model's understanding, capturing more abstract and high-level features important for classification.

5. Flatten Layer:

- Flattens the 2D feature maps into a 1D vector in preparation for the fully connected layers.

6. Fully Connected Layer:

- Dense Layer: 64 units, ReLU activation.
- Dropout Layer: 30% dropout rate.
- This layer learns complex relationships between features, while the Dropout regularizes the model to prevent overfitting.

7. Output Layer:

- Dense Layer: 1 unit, Sigmoid activation.
- Outputs a probability score for binary classification

```
cnn_model_2 = models.Sequential([
    layers.Conv2D(16, (3, 3), activation="relu", padding="same",
input_shape=(image_size, image_size, 3)),
    layers.Conv2D(16, (3, 3), activation="relu", padding="same"),
    layers.MaxPooling2D((2,2)),

    layers.Conv2D(32, (3,3), activation="relu", padding="same"),
    layers.Conv2D(32, (3,3), activation="relu", padding="same"),
    layers.MaxPool2D((2,2)),

    layers.Conv2D(64, (3,3), activation="relu", padding="same"),
    layers.MaxPooling2D((2,2)),

    layers.Flatten(),
    layers.Dense(64, activation="relu"),
    layers.Dropout(0.3),
    layers.Dense(1, activation="sigmoid")
])

cnn_model_2.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0
.005), loss="binary_crossentropy", metrics=["accuracy"])

cnn_model_2.summary()
Model: "sequential_2"
```

Layer (type) Param #	Output Shape	
conv2d_4 (Conv2D) 448	(None, 64, 64, 16)	
conv2d_5 (Conv2D) 2,320	(None, 64, 64, 16)	
max_pooling2d_4 (MaxPooling2D) 0	(None, 32, 32, 16)	
conv2d_6 (Conv2D) 4,640	(None, 32, 32, 32)	
conv2d_7 (Conv2D) 9,248	(None, 32, 32, 32)	
max_pooling2d_5 (MaxPooling2D) 0	(None, 16, 16, 32)	
conv2d_8 (Conv2D) 18,496	(None, 16, 16, 64)	
max_pooling2d_6 (MaxPooling2D) 0	(None, 8, 8, 64)	
flatten_2 (Flatten) 0	(None, 4096)	
dense_4 (Dense) 262,208	(None, 64)	
dropout_2 (Dropout) 0	(None, 64)	
dense_5 (Dense)	(None, 1)	

65 |

Total params: 297,425 (1.13 MB)

Trainable params: 297,425 (1.13 MB)

Non-trainable params: 0 (0.00 B)

As perviously before after creating the CNN model call the fit() method on it to perfrom the training on the modle. The traing will use the training subset and train over 20 epochs with a batch_size of 32 and with a validation split of 20%

```
trained_cnn_model_2 = cnn_model_2.fit(x_train, y_train, epochs=20,  
batch_size=32, validation_split=0.2)
```

Epoch 1/20

4401/4401 _____ 26s 6ms/step - accuracy: 0.5946 - loss: 0.6766 - val_accuracy: 0.6002 - val_loss: 0.6729

Epoch 2/20

4401/4401 _____ 24s 6ms/step - accuracy: 0.5960 - loss: 0.6747 - val_accuracy: 0.6002 - val_loss: 0.6731

Epoch 3/20

4401/4401 _____ 24s 5ms/step - accuracy: 0.5951 - loss: 0.6751 - val_accuracy: 0.6002 - val_loss: 0.6730

Epoch 4/20

4401/4401 _____ 24s 6ms/step - accuracy: 0.5931 - loss: 0.6758 - val_accuracy: 0.6002 - val_loss: 0.6732

Epoch 5/20

4401/4401 _____ 24s 5ms/step - accuracy: 0.5935 - loss: 0.6757 - val_accuracy: 0.6002 - val_loss: 0.6731

Epoch 6/20

4401/4401 _____ 24s 5ms/step - accuracy: 0.5918 - loss: 0.6763 - val_accuracy: 0.6002 - val_loss: 0.6734

Epoch 7/20

4401/4401 _____ 24s 5ms/step - accuracy: 0.5946 - loss: 0.6753 - val_accuracy: 0.6002 - val_loss: 0.6729

Epoch 8/20

4401/4401 _____ 25s 6ms/step - accuracy: 0.5929 - loss: 0.6759 - val_accuracy: 0.6002 - val_loss: 0.6730

Epoch 9/20

4401/4401 _____ 24s 6ms/step - accuracy: 0.5917 - loss: 0.6764 - val_accuracy: 0.6002 - val_loss: 0.6729

Epoch 10/20

4401/4401 _____ 24s 6ms/step - accuracy: 0.5921 - loss: 0.6762 - val_accuracy: 0.6002 - val_loss: 0.6730

Epoch 11/20

4401/4401 _____ 25s 6ms/step - accuracy: 0.5934 - loss: 0.6757 - val_accuracy: 0.6002 - val_loss: 0.6729

```

Epoch 12/20
4401/4401 _____ 25s 6ms/step - accuracy: 0.5944 - loss:
0.6753 - val_accuracy: 0.6002 - val_loss: 0.6731
Epoch 13/20
4401/4401 _____ 24s 6ms/step - accuracy: 0.5917 - loss:
0.6763 - val_accuracy: 0.6002 - val_loss: 0.6730
Epoch 14/20
4401/4401 _____ 24s 5ms/step - accuracy: 0.5922 - loss:
0.6762 - val_accuracy: 0.6002 - val_loss: 0.6729
Epoch 15/20
4401/4401 _____ 24s 5ms/step - accuracy: 0.5951 - loss:
0.6751 - val_accuracy: 0.6002 - val_loss: 0.6730
Epoch 16/20
4401/4401 _____ 24s 5ms/step - accuracy: 0.5951 - loss:
0.6750 - val_accuracy: 0.6002 - val_loss: 0.6730
Epoch 17/20
4401/4401 _____ 24s 6ms/step - accuracy: 0.5970 - loss:
0.6743 - val_accuracy: 0.6002 - val_loss: 0.6730
Epoch 18/20
4401/4401 _____ 24s 6ms/step - accuracy: 0.5925 - loss:
0.6760 - val_accuracy: 0.6002 - val_loss: 0.6730
Epoch 19/20
4401/4401 _____ 24s 6ms/step - accuracy: 0.5923 - loss:
0.6761 - val_accuracy: 0.6002 - val_loss: 0.6729
Epoch 20/20
4401/4401 _____ 24s 6ms/step - accuracy: 0.5933 - loss:
0.6758 - val_accuracy: 0.6002 - val_loss: 0.6730

```

After training the second model the next step is to evaluate the performance of the trained model by using the test subset to score the accuracy of the model. The accuracy of the second model resulted in an accuracy score of about 59.5%. Which is a significantly lower score than the first model. This may be explained by several factors such as the second model utilized a smaller dense layer of 64 units after extracting features, which limited its ability to combine information for final predictions. Also, the high learning rate of 0.005 likely caused unstable updates, which mostly likely prevented the second model from converging properly.

```

cnn_model_2_test_loss, cnn_model_2_test_acc =
cnn_model_2.evaluate(x_test, y_test, verbose=2)


print(f"Test Acc: {cnn_model_2_test_acc:.3f}")

1376/1376 - 3s - 2ms/step - accuracy: 0.5949 - loss: 0.6750
Test Acc: 0.595

```

Plotting the second model's Training vs Validation Accuracy plot visually shows how the validation accuracy did not converge or keep up with the training accuracy. As well as showing how the denser the model did not help it in increasing the accuracy with the number of epochs.

```
plt.plot(trained_cnn_model_2.history["accuracy"], label="Train Acc")
plt.plot(trained_cnn_model_2.history["val_accuracy"], label="Val Acc")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()
plt.title("Model 2: Training vs Validation Accuracy")
plt.show()
```



Model 2 Tuning

In an effort to improve the accuracy of the second model, some tuning will be attempted. The following parameters were tuned in order to increase performance:

- Model depth was decreased to 2 convolutional blocks in an effort to simplify the second model to reduced overfitting risk.
- Dense Layer was increased to 128 units for stronger feature learning.
- The dropout rate was increased to 40% for better regularization.
- Earlystopping was added to stop training when validation loss plateaus.

```
cnn_model_2_tuned = models.Sequential([
    layers.Conv2D(16, (3, 3), activation="relu", padding="same",
input_shape=(64,64,3)),
```

```

layers.Conv2D(16, (3,3), activation="relu", padding="same"),
layers.MaxPooling2D((2,2)),

layers.Conv2D(32, (3,3), activation="relu", padding="same"),
layers.Conv2D(32, (3,3), activation="relu", padding="same"),
layers.MaxPooling2D((2,2)),

layers.Flatten(),
layers.Dense(128, activation="relu"),
layers.Dropout(0.4),
layers.Dense(1, activation="sigmoid")
])

cnn_model_2_tuned.compile(optimizer=tf.keras.optimizers.Adam(learning_
rate=0.005), loss="binary_crossentropy", metrics=["accuracy"])

early_stop = tf.keras.callbacks.EarlyStopping(monitor="val_loss",
patience=5, restore_best_weights=True)

cnn_model_2_tuned.summary()
Model: "sequential_3"

```

Layer (type) Param #	Output Shape	
conv2d_9 (Conv2D) 448	(None, 64, 64, 16)	
conv2d_10 (Conv2D) 2,320	(None, 64, 64, 16)	
max_pooling2d_7 (MaxPooling2D) 0	(None, 32, 32, 16)	
conv2d_11 (Conv2D) 4,640	(None, 32, 32, 32)	
conv2d_12 (Conv2D) 9,248	(None, 32, 32, 32)	
max_pooling2d_8 (MaxPooling2D)	(None, 16, 16, 32)	

0				
		flatten_3 (Flatten)	(None, 8192)	
0				
		dense_6 (Dense)	(None, 128)	
1,048,704				
		dropout_3 (Dropout)	(None, 128)	
0				
		dense_7 (Dense)	(None, 1)	
129				
Total params: 1,065,489 (4.06 MB)				
Trainable params: 1,065,489 (4.06 MB)				
Non-trainable params: 0 (0.00 B)				

After tuning fit the model with the testing subset as previously done.

```
trained_cnn_model_2_tuned = cnn_model_2_tuned.fit(x_train, y_train,
epochs=30, validation_split=0.2, callbacks=[early_stop], verbose=2)

Epoch 1/30
4401/4401 - 22s - 5ms/step - accuracy: 0.5935 - loss: 0.6760 -
val_accuracy: 0.6002 - val_loss: 0.6731
Epoch 2/30
4401/4401 - 21s - 5ms/step - accuracy: 0.5937 - loss: 0.6756 -
val_accuracy: 0.6002 - val_loss: 0.6729
Epoch 3/30
4401/4401 - 21s - 5ms/step - accuracy: 0.5937 - loss: 0.6756 -
val_accuracy: 0.6002 - val_loss: 0.6733
Epoch 4/30
4401/4401 - 21s - 5ms/step - accuracy: 0.5937 - loss: 0.6756 -
val_accuracy: 0.6002 - val_loss: 0.6730
Epoch 5/30
4401/4401 - 22s - 5ms/step - accuracy: 0.5937 - loss: 0.6756 -
val_accuracy: 0.6002 - val_loss: 0.6729
Epoch 6/30
4401/4401 - 22s - 5ms/step - accuracy: 0.5937 - loss: 0.6756 -
val_accuracy: 0.6002 - val_loss: 0.6731
Epoch 7/30
```


```
4401/4401 - 22s - 5ms/step - accuracy: 0.5937 - loss: 0.6756 -  
val_accuracy: 0.6002 - val_loss: 0.6731
```

Again evaluate the tuned second model to check if the tuning helped increase the accuracy score. However the accuracy score stayed the same at 59.5%. This could be due to the tuning changes were not impactful enough for the model to have any increase/decrease in accuracy.

```
cnn_model_2_tuned_test_loss, cnn_model_2_tuned_test_acc =  
cnn_model_2_tuned.evaluate(x_test, y_test, verbose=2)  
  
print(f"Test Acc: {cnn_model_2_tuned_test_acc:.3f}")  
  
1376/1376 - 2s - 2ms/step - accuracy: 0.5949 - loss: 0.6750  
Test Acc: 0.595
```

Plotting the Training vs Validation Accuracy plot shows the same results as the original second model, which shows that the tuned hyperparameters were not impactful enough to change the predictiveness of the model.

```
plt.plot(trained_cnn_model_2_tuned.history["accuracy"], label="Train  
Acc")  
plt.plot(trained_cnn_model_2_tuned.history["val_accuracy"], label="Val  
Acc")  
plt.xlabel("Epoch")  
plt.ylabel("Accuracy")  
plt.legend()  
plt.title("Model 2 Tuned: Training vs Validation Accuracy")  
plt.show()
```





In order to visually see how well each model performed a bar plot of all the model accuracies was created. The bar plot shows that the second CNN model vs tuned model performed exactly the same. The bar chart also shows that the first tuned CNN model performed the best.

```
model_accs = [cnn_model_1_test_acc, cnn_model_1_tuned_test_acc,
cnn_model_2_test_acc, cnn_model_2_tuned_test_acc]

model_names = ["CNN Model 1", "CNN Model 1 Tuned", "CNN Model 2", "CNN
Model 2 Tuned"]

plt.figure(figsize=(10, 6))
plt.bar(model_names, model_accs, edgecolor="black")
plt.ylim(0, 1)
plt.ylabel("Test Accuracy")
plt.title("Comparision of Model Test Accuracies")
plt.grid(axis="y", linestyle="--", alpha=0.7)
plt.xticks(rotation=15)
plt.show()
```



6. Kaggle Results

Now that the models have been tuned and trained that will be used to predict a labels for a seprate traning set of data. Fist, however in order to free up some resources the split training data will be deleted from memory.

```
del x_train, x_test, y_train, y_test
gc.collect()

11594
```

Next a helper fucntion is created to iterate and load all the test TIF images and then create a numpy array that the models can use to predict the labels.

```
def prepare_test_dataset(data_dir:str):
    x_test_new = []
    ids = []

    for file_name in os.listdir(data_dir):
        img_path = os.path.join(data_dir, file_name)
        img = load_img(img_path, target_size=(64,64))
        img_array = img_to_array(img) / 255.0
        x_test_new.append(img_array)
        ids.append(file_name.split(".")[0])
```

```

        return np.array(x_test_new), ids

x_test_new, ids = prepare_test_dataset(test_data_path)

```

The following code blocks will call each model that was trained and use the new testing dataset to predict the labels and create a pandas dataframe.

```

preds_model_1 = cnn_model_1.predict(x_test_new)
pred_labels_model_1 = (preds_model_1 > 0.5).astype(int).flatten()
cnn_model_1_results = pd.DataFrame({"id":ids, "label":
pred_labels_model_1})

```

1796/1796 ————— 3s 2ms/step

```

preds_model_1_tuned = cnn_model_1_tuned.predict(x_test_new)
pred_labels_model_1_tuned = (preds_model_1_tuned >
0.5).astype(int).flatten()
cnn_model_1_tuned_results = pd.DataFrame({"id": ids, "label":
pred_labels_model_1_tuned})

```

1796/1796 ————— 7s 4ms/step

```

preds_model_2 = cnn_model_2.predict(x_test_new)
pred_labels_model_2 = (preds_model_2 > 0.5).astype(int).flatten()
cnn_model_2_results = pd.DataFrame({"id": ids, "label":
pred_labels_model_2})

```

1796/1796 ————— 5s 3ms/step

```

preds_model_2_tuned = cnn_model_2_tuned.predict(x_test_new)
preds_labels_model_2_tuned = (preds_model_2_tuned >
0.5).astype(int).flatten()
cnn_model_2_tuned_results = pd.DataFrame({"id": ids, "label":
preds_labels_model_2_tuned})

```

1796/1796 ————— 4s 2ms/step

After all the models have predicted and created a pandas dataframe the dataframe will be saved locally as a csv. The csv will be submitted to kaggle for scoring.

```

preds_path = os.path.join(os.getcwd(), "model_preds")

cnn_model_1_results.to_csv(os.path.join(preds_path,
"cnn_model_1.csv"), index=False)
cnn_model_1_tuned_results.to_csv(os.path.join(preds_path,
"cnn_model_1_tuned.csv"), index=False)
cnn_model_2_results.to_csv(os.path.join(preds_path,
"cnn_model_2.csv"), index=False)

```

```
cnn_model_2_tuned_results.to_csv(os.path.join(preds_path,
"cnm_model_2_tuned.csv"), index=False)
```

The following screenshot shows the results of each model on Kaggle. As expected the first model that was tuned scored the highest.

Submission and Description	Private Score 	Public Score 	Selected
 cnn_model_2_tuned.csv Complete (after deadline) · now	0.5000	0.5000	☐
 cnn_model_2.csv Complete (after deadline) · 17s ago	0.5000	0.5000	☐
 cnn_model_1_tuned.csv Complete (after deadline) · 29s ago	0.7817	0.8219	☐
 cnn_model_1.csv Complete (after deadline) · 1m ago	0.7773	0.8160	☐

Conclusion

Throughout the development and evaluation of multiple CNN architectures, several important learnings emerged. Simpler models with fewer convolutional layers and moderately sized dense layers, such as the first original model, achieved better performance compared to deeper, more complex models like the second original model. This highlighted that for relatively small datasets, deeper architectures with many convolutional layers may not generalize well without sufficient data and tuning. Tuning efforts, such as increasing the number of filters, adjusting dropout rates, adding EarlyStopping, and enlarging the dense layers, generally helped improve the models' generalization and stability. However, extremely high learning rates, as used in the second model, did not significantly benefit performance and may have hindered convergence. In the future, improvements could include applying data augmentation to artificially expand the training set, using more sophisticated regularization techniques like batch normalization, exploring learning rate scheduling, and experimenting with pre-trained models through transfer learning to leverage previously learned features. Overall, this project reinforced the importance of matching model complexity to dataset size and carefully tuning hyperparameters to optimize model performance.

References

1. [Histopathologic Cancer Detection](#)
2. [Tensorflow Docs](#)
3. [Keras API Docs](#)
4. [Tensorflow CNN Tutorial](#)

Github Repo

https://github.com/charliearvizu-edu/cnn_cancer_detection/tree/dev